

Sridhar Tuli
Assignment 9 Spark

Question:

You can find adult dataset of UCI here (adult.data)

1. Write Python/Java code to cluster the data on at least 3 attributes (for example: age, workclass, education).
2. Execute the above code in Apache Spark.
3. Using spark transformations/actions, find
 - i) The country from which the highest number of adults are there except USA.
 - ii) No of people who are at least Masters and works in "Tech-support".
 - iii) No of unmarried male who works in "Local_govt".

```
adult_analysis.py > ...
1  import pandas as pd
2  from pyspark.sql import SparkSession
3  from pyspark.sql.functions import col
4  from pyspark.ml.feature import VectorAssembler
5  from pyspark.ml.clustering import KMeans
6  from sklearn.preprocessing import LabelEncoder
7
8  spark = SparkSession.builder \
9      .appName("AdultDatasetClustering") \
10     .getOrCreate()
11
12  columns = ["age", "workclass", "fnlwgt", "education", "education_num", "marital_status",
13            "occupation", "relationship", "race", "sex", "capital_gain", "capital_loss",
14            "hours_per_week", "native_country", "income"]
15
16  df = pd.read_csv("adult.data", names=columns, na_values="?")
17
18  label_encoders = {}
19  categorical_columns = ["workclass", "education", "native_country", "occupation", "marital_status", "sex"]
20  for col_name in categorical_columns:
21      le = LabelEncoder()
22      df[col_name] = le.fit_transform(df[col_name].fillna("Unknown"))
23      label_encoders[col_name] = le
24
25  spark_df = spark.createDataFrame(df)
26
27  features = ["age", "workclass", "education"]
28  assembler = VectorAssembler(inputCols=features, outputCol="features")
29  cluster_df = assembler.transform(spark_df)
30
31  kmeans = KMeans(k=3, seed=1)
32  model = kmeans.fit(cluster_df)
33  predictions = model.transform(cluster_df)
34  print("\n\nClustering Results:\n\n")
35  predictions.select("age", "workclass", "education", "prediction").show(5)
36
37  usa_encoded = label_encoders['native_country'].transform([' United-States'])[0]
38
39  highest_country = (spark_df.filter(col("native_country") != usa_encoded)
40                    .groupBy("native_country")
41                    .count()
42                    .orderBy(col("count").desc())
43                    .first())
44
45  if highest_country:
46      highest_country_name = label_encoders['native_country'].inverse_transform([highest_country['native_country']])[0]
47      print(f"\n\nThe country with the highest number of adults except USA: {highest_country_name}\n\n")
```

```

48
49 masters_encoded = label_encoders['education'].transform([' Masters'])[0]
50 doctorate_encoded = label_encoders['education'].transform([' Doctorate'])[0]
51 tech_support_encoded = label_encoders['occupation'].transform([' Tech-support'])[0]
52
53 masters_techsupport = (spark_df.filter(
54     ((col("education") == masters_encoded) | (col("education") == doctorate_encoded)) &
55     (col("occupation") == tech_support_encoded)
56 ).count())
57 print(f"\n\nNumber of people who are at least Masters and work in Tech-support: {masters_techsupport}\n\n")
58
59 local_gov_encoded = label_encoders['workclass'].transform([' Local-gov'])[0]
60 male_encoded = label_encoders['sex'].transform([' Male'])[0]
61
62 married_statuses = [
63     ' Married-civ-spouse',
64     ' Married-spouse-absent',
65     ' Married-AF-spouse'
66 ]
67 married_encoded = [label_encoders['marital_status'].transform([status])[0] for status in married_statuses]
68
69 unmarried_male_localgov = (spark_df.filter(
70     (~col("marital_status").isin(married_encoded)) &
71     (col("sex") == male_encoded) &
72     (col("workclass") == local_gov_encoded)
73 ).count())
74 print(f"\n\nNumber of unmarried males working in Local-gov: {unmarried_male_localgov}\n\n")
75
76 spark.stop()

```

OUTPUT:

Write Python/Java code to cluster the data on at least 3 attributes (for example: age, workclass, education).

24/11/11 10:08:23 INFO CodeGenerator: Code generated in 5.007011 ms

age	workclass	education	prediction
39	7	9	1
50	6	9	2
38	4	11	1
53	4	1	2
28	4	9	0
37	4	12	1
49	4	6	1
52	6	11	2
31	4	12	0
42	4	9	1
37	4	15	1
30	7	9	0
23	4	9	0
32	4	7	0
40	4	8	1
34	4	5	1
25	6	11	0
32	4	11	0
38	4	1	1
43	6	12	1

only showing top 20 rows

3. Using spark transformations/actions, find

i) The country from which the highest number of adults are there except USA.

Corresponding Code:

```
usa_encoded = label_encoders['native_country'].transform([' United-States'])[0]

highest_country = (spark_df.filter(col("native_country") != usa_encoded)
                    .groupBy("native_country")
                    .count()
                    .orderBy(col("count").desc())
                    .first())

if highest_country:
    highest_country_name = label_encoders['native_country'].inverse_transform([highest_country['native_country']])[0]
    print(f"\n\nThe country with the highest number of adults except USA: {highest_country_name}\n\n")

masters_encoded = label_encoders['education'].transform([' Masters'])[0]
doctorate_encoded = label_encoders['education'].transform([' Doctorate'])[0]
tech_support_encoded = label_encoders['occupation'].transform([' Tech-support'])[0]

masters_techsupport = (spark_df.filter(
    ((col("education") == masters_encoded) | (col("education") == doctorate_encoded)) &
    (col("occupation") == tech_support_encoded)
).count())
print(f"\n\nNumber of people who are at least Masters and work in Tech-support: {masters_techsupport}\n\n")

local_gov_encoded = label_encoders['workclass'].transform([' Local-gov'])[0]
male_encoded = label_encoders['sex'].transform([' Male'])[0]
```

Output:

```
24/11/11 10:08:23 INFO DAGScheduler: Job 21 finished: first at /home/solomons/AWS/Spark/adult_analysis.py:43, took 0.041203 s

The country with the highest number of adults except USA: Mexico

24/11/11 10:08:23 INFO CodeGenerator: Code generated in 5.21966 ms
```

ii) No of people who are at least Masters and works in “Tech-support”.

Corresponding Code:

```
masters_encoded = label_encoders['education'].transform([' Masters'])[0]
doctorate_encoded = label_encoders['education'].transform([' Doctorate'])[0]
tech_support_encoded = label_encoders['occupation'].transform([' Tech-support'])[0]

masters_techsupport = (spark_df.filter(
    ((col("education") == masters_encoded) | (col("education") == doctorate_encoded)) &
    (col("occupation") == tech_support_encoded)
).count())
print(f"\n\nNumber of people who are at least Masters and work in Tech-support: {masters_techsupport}\n\n")
```

Output:

```
24/11/11 10:08:23 INFO DAGScheduler: Job 23 finished: count at DirectMethodHandleAccessor.java:103, took 0.016860 s

Number of people who are at least Masters and work in Tech-support: 40

24/11/11 10:08:23 INFO CodeGenerator: Code generated in 7.149224 ms
```

iii) No of unmarried male who works in “Local_govt”.

Corresponding Code:

```
married_statuses = [
    'Married-civ-spouse',
    'Married-spouse-absent',
    'Married-AF-spouse'
]
married_encoded = [label_encoders['marital_status'].transform([status])[0] for status in married_statuses]

unmarried_male_localgov = (spark_df.filter(
    (~col("marital_status").isin(married_encoded)) &
    (col("sex") = male_encoded) &
    (col("workclass") = local_gov_encoded)
).count())
print(f"\n\nNumber of unmarried males working in Local-gov: {unmarried_male_localgov}\n\n")
```

Output:

```
24/11/11 10:08:24 INFO DAGScheduler: Job 25 finished: count at DirectMethodHandleAccessor.java:103, took 0.009877 s

Number of unmarried males working in Local-gov: 384

24/11/11 10:08:24 INFO SparkContext: SparkContext is stopping with exitCode 0.
```

AT SPARK WEBSITE

Spark Jobs (?)

User: solomons
Total Uptime: 12 s
Scheduling Mode: FIFO
Completed Jobs: 26

▶ Event Timeline

▼ Completed Jobs (26)

Page: 1

1 Pages. Jump to 1. Show 100 items in a page. Go

Job id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
25	count at DirectMethodHandleAccessor.java:103 count at DirectMethodHandleAccessor.java:103	2024/11/11 10:11:05	8 ms	1/1 (1 skipped)	1/1 (12 skipped)
24	count at DirectMethodHandleAccessor.java:103 count at DirectMethodHandleAccessor.java:103	2024/11/11 10:11:05	0.2 s	1/1	12/12
23	count at DirectMethodHandleAccessor.java:103 count at DirectMethodHandleAccessor.java:103	2024/11/11 10:11:05	12 ms	1/1 (1 skipped)	1/1 (12 skipped)
22	count at DirectMethodHandleAccessor.java:103 count at DirectMethodHandleAccessor.java:103	2024/11/11 10:11:05	0.2 s	1/1	12/12
21	first at /home/solomons/AWS/Spark/adult_analysis.py:43 first at /home/solomons/AWS/Spark/adult_analysis.py:43	2024/11/11 10:11:05	32 ms	1/1 (1 skipped)	1/1 (12 skipped)
20	first at /home/solomons/AWS/Spark/adult_analysis.py:43 first at /home/solomons/AWS/Spark/adult_analysis.py:43	2024/11/11 10:11:05	0.2 s	1/1	12/12

We are getting the information about the task happening
My laptop has 6 Cores, therefore 12 threads
Thats why it’s showing 1/12 in the image, telling that 1 task has been completed out of 12
parallely running ones.

=====END OF FILE=====